

```
import pandas as pd
import numpy as np
import torch
import torchvision
import random
from PIL import Image
import cv2
import os
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
!pip install evaluate jiwer
```

Collecting evaluate

Downloading evaluate-0.4.0-py3-none-any.whl (81 kB)

81.4/81.4 kB 607.4 kB/s eta

0:00:00a 0:00:01

Requirement already satisfied: tqdm<4.62.1 in /opt/conda/lib/python3.7/site-packages (from evaluate) (4.64.0)

Requirement already satisfied: xxhash in

/opt/conda/lib/python3.7/site-packages (from evaluate) (3.0.0)

Requirement already satisfied: fsspec[http]>=2021.05.0 in

/opt/conda/lib/python3.7/site-packages (from evaluate) (2022.8.2)

Requirement already satisfied: pandas in

/opt/conda/lib/python3.7/site-packages (from evaluate) (1.3.5)

Requirement already satisfied: multiprocessing in

/opt/conda/lib/python3.7/site-packages (from evaluate) (0.70.13)

Requirement already satisfied: importlib-metadata in

/opt/conda/lib/python3.7/site-packages (from evaluate) (4.13.0)

Requirement already satisfied: packaging in

/opt/conda/lib/python3.7/site-packages (from evaluate) (21.3)

Requirement already satisfied: datasets>=2.0.0 in

/opt/conda/lib/python3.7/site-packages (from evaluate) (2.1.0)

Requirement already satisfied: dill in /opt/conda/lib/python3.7/site-packages (from evaluate) (0.3.5.1)

Requirement already satisfied: responses<0.19 in

/opt/conda/lib/python3.7/site-packages (from evaluate) (0.18.0)

Requirement already satisfied: huggingface-hub>=0.7.0 in

/opt/conda/lib/python3.7/site-packages (from evaluate) (0.10.1)

Requirement already satisfied: numpy>=1.17 in

/opt/conda/lib/python3.7/site-packages (from evaluate) (1.21.6)

Requirement already satisfied: requests>=2.19.0 in

/opt/conda/lib/python3.7/site-packages (from evaluate) (2.28.1)

Collecting levenshtein==0.20.2

Downloading Levenshtein-0.20.2-cp37-cp37m-

manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.4 MB)

1.4/1.4 MB 4.4 MB/s eta

0:00:0000:0100:01

Requirement already satisfied: rapidfuzz<3.0.0,>=2.3.0 in

/opt/conda/lib/python3.7/site-packages (from levenshtein==0.20.2->jiwer) (2.11.1)

Requirement already satisfied: aiohttp in

/opt/conda/lib/python3.7/site-packages (from datasets>=2.0.0->evaluate) (3.8.1)
Requirement already satisfied: pyarrow>=5.0.0 in /opt/conda/lib/python3.7/site-packages (from datasets>=2.0.0->evaluate) (5.0.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /opt/conda/lib/python3.7/site-packages (from huggingface-hub>=0.7.0->evaluate) (4.1.1)
Requirement already satisfied: pyyaml>=5.1 in /opt/conda/lib/python3.7/site-packages (from huggingface-hub>=0.7.0->evaluate) (6.0)
Requirement already satisfied: filelock in /opt/conda/lib/python3.7/site-packages (from huggingface-hub>=0.7.0->evaluate) (3.7.1)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /opt/conda/lib/python3.7/site-packages (from packaging->evaluate) (3.0.9)
Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.7/site-packages (from requests>=2.19.0->evaluate) (3.3)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /opt/conda/lib/python3.7/site-packages (from requests>=2.19.0->evaluate) (1.26.12)
Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.7/site-packages (from requests>=2.19.0->evaluate) (2022.9.24)
Requirement already satisfied: charset-normalizer<3,>=2 in /opt/conda/lib/python3.7/site-packages (from requests>=2.19.0->evaluate) (2.1.0)
Requirement already satisfied: zipp>=0.5 in /opt/conda/lib/python3.7/site-packages (from importlib-metadata->evaluate) (3.8.0)
Requirement already satisfied: python-dateutil>=2.7.3 in /opt/conda/lib/python3.7/site-packages (from pandas->evaluate) (2.8.2)
Requirement already satisfied: pytz>=2017.3 in /opt/conda/lib/python3.7/site-packages (from pandas->evaluate) (2022.1)
Requirement already satisfied: aiosignal>=1.1.2 in /opt/conda/lib/python3.7/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (1.2.0)
Requirement already satisfied: asyncctest==0.13.0 in /opt/conda/lib/python3.7/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (0.13.0)
Requirement already satisfied: frozenlist>=1.1.1 in /opt/conda/lib/python3.7/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (1.3.0)
Requirement already satisfied: yarl<2.0,>=1.0 in /opt/conda/lib/python3.7/site-packages (from aiohttp->datasets>=2.0.0->evaluate) (1.7.2)
Requirement already satisfied: attrs>=17.3.0 in

```

/opt/conda/lib/python3.7/site-packages (from aiohttp->datasets>=2.0.0-
>evaluate) (21.4.0)
Requirement already satisfied: async-timeout<5.0,>=4.0.0a3 in
/opt/conda/lib/python3.7/site-packages (from aiohttp->datasets>=2.0.0-
>evaluate) (4.0.2)
Requirement already satisfied: multidict<7.0,>=4.5 in
/opt/conda/lib/python3.7/site-packages (from aiohttp->datasets>=2.0.0-
>evaluate) (6.0.2)
Requirement already satisfied: six>=1.5 in
/opt/conda/lib/python3.7/site-packages (from python-dateutil>=2.7.3-
>pandas->evaluate) (1.15.0)
Installing collected packages: levenshtein, jiwer, evaluate
  Attempting uninstall: levenshtein
    Found existing installation: Levenshtein 0.20.7
    Uninstalling Levenshtein-0.20.7:
      Successfully uninstalled Levenshtein-0.20.7
ERROR: pip's dependency resolver does not currently take into account
all the packages that are installed. This behaviour is the source of
the following dependency conflicts.
python-levenshtein 0.20.7 requires Levenshtein==0.20.7, but you have
levenshtein 0.20.2 which is incompatible.
Successfully installed evaluate-0.4.0 jiwer-2.5.1 levenshtein-0.20.2
WARNING: Running pip as the 'root' user can result in broken
permissions and conflicting behaviour with the system package manager.
It is recommended to use a virtual environment instead:
https://pip.pypa.io/warnings/venv

```

```

"""Seed everything!"""
random.seed(42)
os.environ['PYTHONHASHSEED'] = str(42)
np.random.seed(42)
torch.manual_seed(42)
torch.cuda.manual_seed(42)
torch.backends.cudnn.deterministic = True
torch.backends.cudnn.benchmark = True
device

```

Load dataset

Tokenize car numbers

Get the list of car numbers

```

def exec_text(path):
    return path[path.find('-') + 1:path.find('.')]

```

```

input_dir_train = '/kaggle/input/labtinkoff/CCPD2019-d11/train'

```

```

car_numbers = [exec_text(path) for path in
os.listdir(input_dir_train)]

```

```

# Get the alphabet of symbols from all car numbers
seq = ''
for car_number in car_numbers:
    seq += car_number
alphabet = ''
for symbol in sorted(set(seq)):
    alphabet += symbol
alphabet

'0123456789ABCDEFGHIJKLMNOPQRSTUVWXYZ 云京冀吉宁川新晋桂沪津浙渝湘琼甘皖粤苏蒙
藏豫贵赣辽鄂闽陕青鲁黑'

OOV_TOKEN = '<OOV>' # out of vocabulary token
CTC_BLANK = '<BLANK>' # token for ctc matrix
PAD_TOKEN = '<PAD>' # padding token

def get_char_map(alphabet):
    """Make from string alphabet character2int dict.
    Add BLANK char for CTC loss and OOV char for out of vocabulary
    symbols."""
    char_map = {value: idx + 3 for (idx, value) in
enumerate(alphabet)}
    char_map[CTC_BLANK] = 0
    char_map[OOV_TOKEN] = 1
    char_map[PAD_TOKEN] = 2
    return char_map

class Tokenizer:
    """Class for encoding and decoding string word to sequence of int
    (and vice versa) using alphabet."""

    def __init__(self, alphabet):
        self.char_map = get_char_map(alphabet)
        self.rev_char_map = {val: key for key, val in
self.char_map.items()}

    def encode(self, word_list):
        enc_words = []
        for word in word_list:
            enc_words.append(
                [self.char_map[char] if char in self.char_map
                 else self.char_map[OOV_TOKEN]
                 for char in word]
            )
        return enc_words

    def get_num_chars(self):
        return len(self.char_map)

```

```

def decode(self, enc_word_list):
    dec_words = []
    for word in enc_word_list:
        word_chars = ''
        for idx, char_enc in enumerate(word):
            if (
                char_enc != self.char_map[OOV_TOKEN]
                and char_enc != self.char_map[CTC_BLANK]
                and not (idx > 0 and char_enc == word[idx - 1])
            ):
                word_chars += self.rev_char_map[char_enc]
        dec_words.append(word_chars)
    return dec_words

tokenizer = Tokenizer(alphabet)

class Laba_dataset(torch.utils.data.Dataset):
    def __init__(self, root, tokenizer, transform=None):
        self.root = root
        self.transform = transform
        self.tokenizer = tokenizer
        self.img_paths = [os.path.join(self.root, img_path) for
img_path in os.listdir(self.root)]
        self.text = [exec_text(path) for path in
os.listdir(self.root)]
        self.enc_text = self.tokenizer.encode(self.text)

    def __getitem__(self, ind):
        img = Image.open(self.img_paths[ind]) # resize
        if self.transform is not None:
            img = self.transform(img) # make some augmentations
        # return image, encoded_text, source_text
        return (img, torch.LongTensor(self.enc_text[ind]),
self.text[ind])

    def __len__(self):
        return len(self.img_paths)

def collate_fn(batch):
    images, enc_texts, texts = zip(*batch)
    images = torch.stack(images, 0)
    enc_pad_texts = torch.nn.utils.rnn.pad_sequence(enc_texts,
batch_first=True, padding_value=tokenizer.char_map[PAD_TOKEN])
    return images, enc_pad_texts, texts

from sklearn.model_selection import train_test_split
batch_size = 128
# some augmentations for model regularization
transform = torchvision.transforms.Compose([
    torchvision.transforms.Resize((32, 128)),

```

```

        torchvision.transforms.RandomRotation(5),
        torchvision.transforms.ColorJitter(),
        torchvision.transforms.GaussianBlur(3),
        torchvision.transforms.ToTensor()
    ])
dataset_full = Laba_dataset(input_dir_train, tokenizer,
transform=transform)
# split full dataset
train_idx, valid_idx =
train_test_split(list(range(len(dataset_full))), train_size=0.9)
dataset = {
    'train': torch.utils.data.Subset(dataset_full, train_idx),
    'valid': torch.utils.data.Subset(dataset_full, valid_idx)
}

dataset_size = {ds: len(dataset[ds]) for ds in ['train', 'valid']}

dataloader = {
    'train': torch.utils.data.DataLoader(
        dataset=dataset['train'], batch_size=batch_size, shuffle=True,
        collate_fn=collate_fn
    ),
    'valid': torch.utils.data.DataLoader(
        dataset=dataset['valid'], batch_size=batch_size,
        shuffle=False, collate_fn=collate_fn
    ),
}

input_dir_test = '/kaggle/input/labtinkoff/CCPD2019-dl1/test'
batch_size = 64
transform_test = torchvision.transforms.Compose([
    torchvision.transforms.Resize((32, 128)),
    torchvision.transforms.ToTensor()
])
dataset_test = Laba_dataset(input_dir_test, tokenizer,
transform=transform_test)
dataloader_test = torch.utils.data.DataLoader(
    dataset=dataset_test, batch_size=batch_size, shuffle=False,
    collate_fn=collate_fn
)

next(iter(dataloader['train']))[0].shape

torch.Size([128, 3, 32, 128])

img = torchvision.transforms.ToPILImage()(dataset_full[173]
[0].squeeze(0)) # take a look to random image
img

```



```
dataset_test[122] # take a look to element from dataset
(tensor([[[[0.5451, 0.5529, 0.5608, ..., 0.4157, 0.4118, 0.4706],
           [0.4431, 0.4039, 0.4039, ..., 0.4157, 0.5020, 0.5647],
           [0.3333, 0.2980, 0.3059, ..., 0.2353, 0.3137, 0.4353],
           ...,
           [0.3333, 0.3137, 0.2667, ..., 0.1608, 0.1490, 0.1843],
           [0.3216, 0.3176, 0.3098, ..., 0.4471, 0.4706, 0.5020],
           [0.3255, 0.3333, 0.3529, ..., 0.3922, 0.4314, 0.4471]]],

         [[0.5059, 0.5059, 0.5059, ..., 0.4745, 0.4706, 0.5373],
          [0.4275, 0.3686, 0.3608, ..., 0.5216, 0.6000, 0.6627],
          [0.3569, 0.3020, 0.2980, ..., 0.3804, 0.4431, 0.5490],
          ...,
          [0.4431, 0.4235, 0.3725, ..., 0.2353, 0.2392, 0.2824],
          [0.4157, 0.4157, 0.4078, ..., 0.5255, 0.5647, 0.6000],
          [0.4157, 0.4235, 0.4431, ..., 0.4784, 0.5255, 0.5412]]],

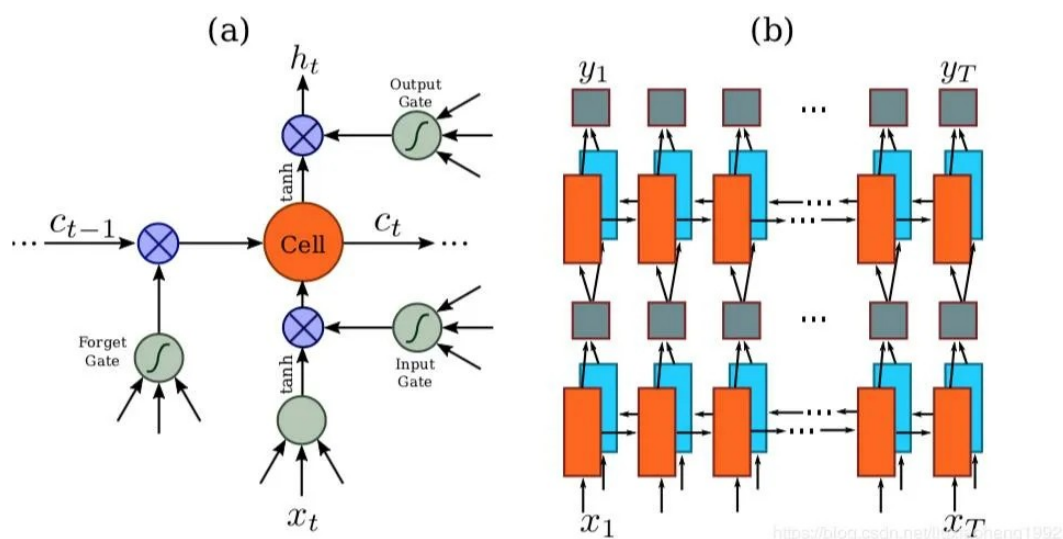
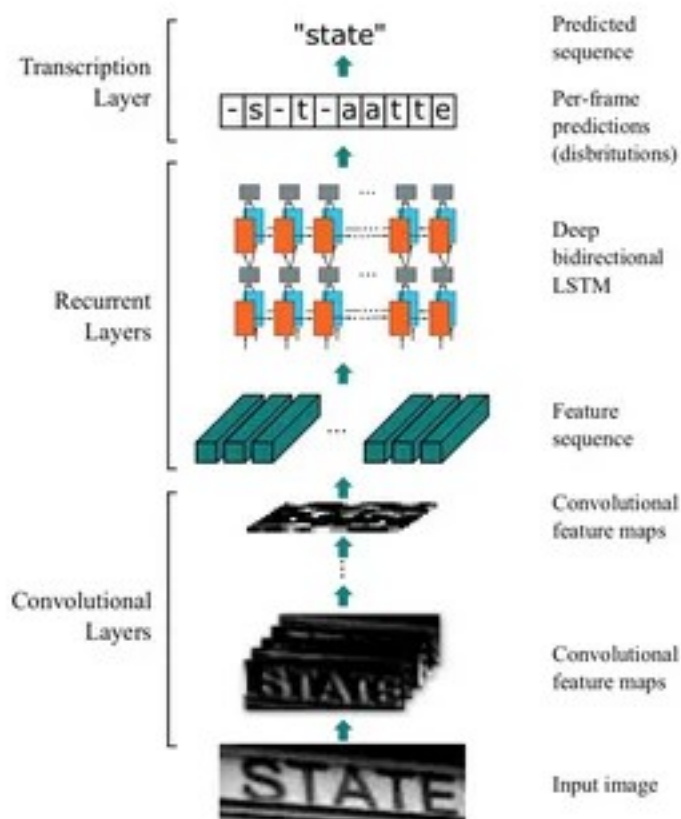
         [[0.6275, 0.6353, 0.6431, ..., 0.5451, 0.5176, 0.5333],
          [0.5961, 0.5765, 0.6000, ..., 0.6000, 0.6353, 0.6471],
          [0.5490, 0.5725, 0.6471, ..., 0.5059, 0.4902, 0.5294],
          ...,
          [0.4314, 0.4471, 0.4510, ..., 0.2980, 0.2039, 0.2000],
          [0.3882, 0.3922, 0.3922, ..., 0.5529, 0.5216, 0.5294],
          [0.3922, 0.3961, 0.4118, ..., 0.4784, 0.4863, 0.4941]]]),
  tensor([54, 13, 5, 8, 33, 8, 12]),
  '皖A25V59')
```

Define model

```
from torch import nn
```

To solve the problem i use CRNN structure

Approximate model structure and LSTM structure



```
class ResNetBlock(nn.Module):
    def __init__(self, in_channels, out_channels, kernel_size=3,
        stride=1, padding=0, dropout=0.15):
        super().__init__()
        self.conv = nn.Conv2d(in_channels, out_channels, kernel_size,
            stride, padding, bias=False)
        self.bn = nn.BatchNorm2d(out_channels)
```



```

        self.relu = nn.LeakyReLU()
        self.dropout = nn.Dropout(dropout)
        self.downsample = None
        if in_channels != out_channels:
            self.downsample = nn.Conv2d(in_channels, out_channels, 1,
stride=2)

    def forward(self, x, identity=True):
        out = self.dropout(self.bn(self.conv(x)))
        if identity:
            if self.downsample is not None:
                x = self.downsample(x)
            return self.relu(out + x)
        else:
            return self.relu(out)

class CNN(nn.Module):
    def __init__(self, in_channels=1, num_layers=2, dropout=0.1):
        super().__init__()
        """
        As feature extractor i use resnet, passing through the cut the
        images are
        transformed from the dimension tensor (C: 1, W: 128, H: 32) to
        the
        dimension tensor (C: 1, W: 4, H: 1)
        """
        self.start = ResNetBlock(3, 64, 7, 1, 0, 0.0)
        self.maxpool = nn.MaxPool2d(3, 2, 1)
        self.blocks1 = nn.ModuleList([ResNetBlock(64, 64, padding=1)
for _ in range(num_layers)])
        self.blocks2 = nn.ModuleList([ResNetBlock(64, 128, padding=1,
stride=2)] + [ResNetBlock(128, 128, padding=1) for _ in
range(num_layers)])
        self.blocks3 = nn.ModuleList([ResNetBlock(128, 256, padding=1,
stride=2)] + [ResNetBlock(256, 256, padding=1) for _ in
range(num_layers)])
        self.blocks4 = nn.ModuleList([ResNetBlock(256, 512, padding=1,
stride=2)] + [ResNetBlock(512, 512, padding=1) for _ in
range(num_layers)])
        self.blocks5 = nn.ModuleList([ResNetBlock(512, 1024,
padding=1, stride=2)] + [ResNetBlock(1024, 1024, padding=1) for _ in
range(num_layers)])
        self.blocks = [self.blocks1, self.blocks2, self.blocks3,
self.blocks4, self.blocks5]

    def forward(self, x):
        out = self.maxpool(self.start(x, identity=False))
        for blocks in self.blocks:
            for block in blocks:
                out = block(out)

```

```

        return out

class BiLSTM(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers,
dropout=0.1):
        super().__init__()
        self.lstm = nn.LSTM(
            input_size, hidden_size, num_layers,
            dropout=dropout, batch_first=True, bidirectional=True)

    def forward(self, x):
        out, _ = self.lstm(x)
        return out

"""
5 ResNet blocks, 3 BiLSTM, Linear Classifier
in order for the model to be able to identify more complex
dependencies in the data, since the dataset allows you to enter
a deep neural network
"""

class CRNN(nn.Module):
    def __init__(
        self, number_class_symbols, time_feature_count=256,
lstm_hidden=256,
        lstm_len=3,
    ):
        super().__init__()
        self.feature_extractor = CNN(dropout=0.15)
        self.avg_pool = nn.AdaptiveAvgPool2d(
            (time_feature_count, time_feature_count))
        self.bilstm = BiLSTM(time_feature_count, lstm_hidden,
lstm_len, dropout=0.15)
        self.classifier = nn.Sequential(
            nn.Linear(lstm_hidden * 2, time_feature_count),
            nn.GELU(),
            nn.Dropout(0.15),
            nn.Linear(time_feature_count, number_class_symbols) # the
model predicts the probability of characters from the alphabet
        )

    def forward(self, x):
        x = self.feature_extractor(x)
        b, c, h, w = x.size()
        x = x.view(b, c * h, w)
        x = self.avg_pool(x)
        x = x.transpose(1, 2)
        x = self.bilstm(x)
        x = self.classifier(x)

```

```

x = nn.functional.log_softmax(x, dim=2).permute(1, 0, 2)
return x

```

Define accuracy metric for evaluate validation dataset

```

class AverageMeter:
    def __init__(self):
        self.reset()

    def reset(self):
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.sum += val * n
        self.count += n
        self.avg = self.sum / self.count

    def get_accuracy(y_true, y_pred):
        scores = []
        for true, pred in zip(y_true, y_pred):
            scores.append(true == pred)
        avg_score = np.mean(scores)
        return avg_score

```

Training loop

functions for saving models

```
import pickle as pkl
```

```

def safe(obj, filename):
    with open(filename, 'wb') as outp:
        pkl.dump(obj, outp)

def read(filename):
    with open(filename, 'rb') as inp:
        return pkl.load(inp)

def weights_init(m):
    classname = m.__class__.__name__
    if type(m) in [nn.Linear, nn.Conv2d, nn.Conv1d]:
        torch.nn.init.xavier_uniform_(m.weight)
        if m.bias is not None:
            m.bias.data.fill_(0.01)
    elif classname.find('BatchNorm') != -1:
        m.weight.data.normal_(1.0, 0.02)
        m.bias.data.fill_(0)

def val_loop(data_loader, model, tokenizer, device):
    acc_avg = AverageMeter()

```

```

    for images, enc_texts, texts in data_loader:
        batch_size = len(texts)
        text_preds = predict(images, model, tokenizer, device)
        acc_avg.update(get_accuracy(texts, text_preds), batch_size)
    print(f'Validation, acc: {acc_avg.avg:.4f}')
    return acc_avg.avg

def predict(images, model, tokenizer, device):
    model.eval()
    images = images.to(device)
    with torch.no_grad():
        output = model(images)
    pred = torch.argmax(output.detach().cpu(), -1).permute(1,
0).numpy()
    text_preds = tokenizer.decode(pred)
    return text_preds

def train_loop(data_loader, model, criterion, optimizer, epoch):
    loss_avg = AverageMeter()
    model.train()
    for images, enc_texts, texts in data_loader:
        model.zero_grad()
        images = images.to(device)
        batch_size = len(texts)
        output = model(images)
        output_lengths = torch.full(
            size=(output.size(1),),
            fill_value=output.size(0),
            dtype=torch.long
        )
        text_lens = torch.LongTensor([len(text) for text in texts]) #
for CTC-loss
        loss = criterion(output, enc_texts, output_lengths, text_lens)
        loss_avg.update(loss.item(), batch_size)
        loss.backward()
        torch.nn.utils.clip_grad_norm_(model.parameters(), 2)
        optimizer.step()
    for param_group in optimizer.param_groups:
        lr = param_group['lr']
    print(f'\nEpoch {epoch}, Loss: {loss_avg.avg:.5f}, LR: {lr:.7f}')
    return loss_avg.avg

accs = {}
def train(dataloader, epochs, trained_model=None,
trained_model_epochs=0):
    train_loader, val_loader = dataloader['train'],
dataloader['valid']
    if trained_model == None:
        model = CRNN(number_class_symbols=tokenizer.get_num_chars())
        model.apply(weights_init)

```

```

        model.to(device)
    else:
        model = trained_model

    criterion = torch.nn.CTCLoss(blank=0, reduction='mean',
zero_infinity=True)
    optimizer = torch.optim.AdamW(model.parameters(), lr=0.001,
                                weight_decay=0.01)
    scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
        optimizer=optimizer, mode='max', factor=0.5, patience=5)
    best_acc = -np.inf
    acc_avg = val_loop(val_loader, model, tokenizer, device)
    for epoch in range(epochs):
        loss_avg = train_loop(train_loader, model, criterion,
optimizer, epoch)
        acc_avg = val_loop(val_loader, model, tokenizer, device)
        accs[acc_avg] = epoch
        scheduler.step(acc_avg)
        if acc_avg > best_acc:
            best_acc = acc_avg
        if trained_model_epochs == 0:
            save(model, f'model_{epoch}')
        else:
            save(model, f'model_{epoch + trained_model_epochs}')

```

train(dataloader, 8)

Validation, acc: 0.0000

Epoch 0, Loss: 1.68883, LR: 0.0010000

Validation, acc: 0.3340

Epoch 1, Loss: 0.23575, LR: 0.0010000

Validation, acc: 0.6538

Epoch 2, Loss: 0.10533, LR: 0.0010000

Validation, acc: 0.9317

Epoch 3, Loss: 0.05218, LR: 0.0010000

Validation, acc: 0.9395

Epoch 4, Loss: 0.04142, LR: 0.0010000

Validation, acc: 0.9573

Epoch 5, Loss: 0.03328, LR: 0.0010000

Validation, acc: 0.9699

Epoch 6, Loss: 0.02488, LR: 0.0010000

Validation, acc: 0.9650

```
Epoch 7, Loss: 0.02500, LR: 0.0010000  
Validation, acc: 0.9668  
  
model = read(f'/kaggle/working/model_{accs[max(accs.keys())]}') # load  
model with best acc on validation  
  
train(dataloader, 10, model, 8) # more epochs  
  
Validation, acc: 0.9684  
  
Epoch 0, Loss: 0.02766, LR: 0.0010000  
Validation, acc: 0.9715  
  
Epoch 1, Loss: 0.02236, LR: 0.0010000  
Validation, acc: 0.9669  
  
Epoch 2, Loss: 0.01953, LR: 0.0010000  
Validation, acc: 0.9771  
  
Epoch 3, Loss: 0.01721, LR: 0.0010000  
Validation, acc: 0.9784  
  
Epoch 4, Loss: 0.01611, LR: 0.0010000  
Validation, acc: 0.9821  
  
Epoch 5, Loss: 0.01570, LR: 0.0010000  
Validation, acc: 0.9808  
  
Epoch 6, Loss: 0.01318, LR: 0.0010000  
Validation, acc: 0.9733  
  
Epoch 7, Loss: 0.01456, LR: 0.0010000  
Validation, acc: 0.9831  
  
Epoch 8, Loss: 0.01173, LR: 0.0010000  
Validation, acc: 0.9848  
  
Epoch 9, Loss: 0.01146, LR: 0.0010000  
Validation, acc: 0.9827  
  
img, enc_label, label = dataset_full[2001]  
  
model = read(f'/kaggle/working/model_{16}') # take model with the best  
acc on validation  
  
pred = predict(img.unsqueeze(0).to(device), model, tokenizer, device)  
# sample pred  
pred  
  
['皖KLJ029']
```

```
real_img = torchvision.transforms.ToPILImage()(img)
real_img
```



Compute metrics

$$\text{CER} = (S + D + I) / N = (S + D + I) / (S + D + C)$$

where

S is the number of substitutions, D is the number of deletions, I is the number of insertions, C is the number of correct characters, N is the number of characters in the reference (N=S+D+C).

```
from evaluate import load
cer = load("cer")

references = dataset_test.text

predictions = []
for imgs, enc_text, text in dataloader_test:
    predictions += predict(imgs, model, tokenizer, device)

cer.compute(predictions=predictions, references=references)

0.0058005800580058
```

CER's output is not always a number between 0 and 1, in particular when there is a high number of insertions. This value is often associated to the percentage of characters that were incorrectly predicted. The lower the value, the better the performance of the ASR system with a CER of 0 being a perfect score.

```
len(references) == len(predictions)
```

True

```
errors = {} # dict of errors {predictions: references}
for pred, refer in zip(predictions, references):
    if cer.compute(predictions=[pred], references=[refer]) != 0.0:
        errors[pred] = refer
```

errors

```
{ '皖 HL108D': '皖 NL108D',
  '皖 AE0082': '皖 AF888S',
  '皖 AN5R29': '冀 AN5R29',
  '皖 AE8F76': '皖 AF8F76',
  '冀 JS2363': '川 JS2363',
  '皖 A06Z16': '皖 AD6Z16',
```

'皖AL1D76': '皖AL1D70',
'豫ADL439': '甘ADL439',
'皖AZZ2D8': '皖AZZ208',
'鄂E6V671': '赣E6V671',
'鲁CKB654': '粤CKB654',
'赣E269JY': '湘E269JY',
'鄂AN08Q3': '豫AN08Q3',
'皖SC716': '皖SCC716',
'皖AB9017': '皖AP901T',
'皖AH1179': '皖AH1178',
'皖AHA18C': '皖AHA180',
'鄂A607MG': '黑A607MG',
'皖A9C714': '皖ADX714',
'苏N71031': '沪N71031',
'闽BB1188': '鄂BB1188',
'皖A75156': '皖A751S6',
'皖A280D5': '皖A28G05',
'皖AL13DD': '皖AL130D',
'皖AZW563': '皖AZW063',
'皖AC38DM': '皖AC380M',
'浙GLQ029': '津GLQ029',
'皖E659B8': '赣E659B8',
'冀FMK521': '鄂FMK521',
'皖AG3327': '皖A03327',
'皖AEDN66': '皖AE0N66',
'皖APX498': '皖ADX498',
'浙CN060V': '闽CN060V',
'皖AN866G': '皖AN866B',
'京Q6099K': '鲁Q6099K',
'苏E160C0': '苏E16GL0',
'苏A68R93': '苏A68RX3',
'皖A050C0': '皖A05000',
'皖MQE698': '皖NQE698',
'京NRB208': '鲁NRB208',
'皖EXX665': '陕EXX665',
'晋BQA986': '鲁BQA986',
'皖NBB293': '粤NBB293',
'皖DR22P': '皖RL222P',
'豫SC6Q25': '粤SC6Q25',
'皖L35P91': '粤L35P91',
'鄂FLK476': '贵FLK476',
'皖AJ2928': '皖AV2928',
'浙DK8733': '沪DK8733',
'粤BL1507': '琼BL1507',
'皖AA79D6': '皖AA7S06',
'皖AP1P17': '皖AP1P13',
'皖A812A5': '皖A81245',
'粤MS8230': '黑MS8230',
'皖HYA266': '皖EYA266',

'浙 EXQ3C3' : '浙 EXQ303',
'京 P76N06' : '京 P76M06',
'皖 AJ6408' : '皖 AKL408',
'皖 ABT220' : '皖 ABT820',
'皖 AZ33DT' : '皖 AZ330T',
'皖 A2M2K8' : '皖 A252K8',
'皖 A6S33D' : '皖 A6S330',
'冀 FKK507' : '鲁 FKK507',
'鄂 A52NH4' : '鄂 A52VH4',
'皖 AB1Q30' : '皖 AB1930',
'浙 B58K96' : '浙 B58KX6',
'皖 A522B0' : '皖 A522R0',
'皖 ANY863' : '皖 ANY833',
'闽 CL7029' : '鄂 CL7029',
'皖 AH3099' : '皖 AH309S',
'皖 ADK195' : '皖 A0K195',
'沪 A0X007' : '蒙 A0X007',
'皖 A3151H' : '云 A3151H',
'鄂 A71XY8' : '皖 A771X8',
'皖 A235H0' : '皖 A235U0',
'豫 LLD156' : '鄂 LLD155',
'皖 AD0926' : '皖 A00926',
'皖 BZ996J' : '皖 RZ996J',
'皖 AD0T99' : '皖 AD0T89',
'皖 G99C66' : '皖 Q99066',
'皖 AB011M' : '皖 AD011M',
'皖 AFQ958' : '云 AFQ958',
'皖 A9Z5D5' : '皖 A9Z505',
'渝 DZ8121' : '晋 DZ8121',
'皖 AD76DD' : '皖 AD760D',
'皖 A18RQ0' : '皖 A18R00',
'粤 SG6U10' : '豫 SG6U10',
'皖 A6X91' : '皖 AX6X91',
'粤 L3N780' : '鄂 L3N780',
'浙 F7H069' : '湘 F7H069',
'赣 M35612' : '沪 M35612',
'皖 AD1P09' : '皖 AP1P09',
'皖 A75W26' : '皖 A75W76',
'皖 AF0A61' : '皖 AP0A61',
'皖 AK088N' : '皖 AK098N',
'皖 FXH016' : '皖 EXH016',
'皖 A368SP' : '浙 A368SP',
'粤 X937L9' : '鲁 Y037L9',
'皖 AK801V' : '皖 AK801W',
'赣 G2L260' : '豫 G2L260',
'皖 A65019' : '皖 A650T9',
'皖 AUT949' : '皖 A0T949',
'皖 AY6A32' : '皖 A96A39',
'皖 AU3Y97' : '冀 AU3Y97',

'皖 AN6255' : '新 AN6255',
'皖 HPS138' : '皖 HPS178',
'皖 AG961C' : '皖 AG961L',
'浙 BSX779' : '沪 BSX779',
'鲁 HFP512' : '浙 HFP512',
'皖 A8Y564' : '皖 A0Y064',
'皖 AWU097' : '皖 AWE097',
'浙 A7304H' : '湘 A7304H',
'豫 A883T8' : '蒙 A883T8',
'皖 AHBV95' : '皖 AH8V95',
'皖 A25605' : '皖 A2W003',
'皖 A829X9' : '皖 A829Y9',
'浙 EHJ267' : '粤 EHJ267',
'皖 AF0K08' : '皖 AF0V08',
'鄂 A02SC9' : '鄂 A02SU9',
'皖 A12376' : '皖 A123X5',
'皖 ASF6C4' : '皖 ASF604',
'晋 C1G999' : '浙 C1G999',
'皖 A918B1' : '皖 A918Q1',
'皖 ACS380' : '皖 AES380',
'鄂 A7829P' : '豫 A7829P',
'湘 A5K75S' : '豫 A5K75S',
'皖 AMS995' : '新 AMS995',
'皖 AMP669' : '皖 AMP600',
'皖 A6915C' : '皖 AJ915C',
'皖 C23D77' : '皖 C23077',
'皖 AD289X' : '皖 AD289W',
'皖 AC3J66' : '皖 AC3J88',
'皖 AVE081' : '皖 AVL084',
'皖 AR316B' : '皖 AR316D',
'京 HA8A15' : '苏 FA8A15',
'鄂 GK3511' : '赣 GK3511',
'皖 AWQ94F' : '皖 AW094F',
'苏 E67N97' : '苏 E67NX7',
'苏 B20QQB' : '苏 B200QB',
'皖 AH81Q2' : '辽 AH81Q2',
'皖 AHE220' : '皖 AHF220',
'皖 AWB8D4' : '皖 AWB804',
'皖 AGS39P' : '皖 AGS308',
'京 PN8230' : '辽 PN8230',
'粤 HA149Y' : '粤 YA149Y',
'皖 A981SU' : '皖 A981S0',
'浙 BB906Y' : '苏 BB906Y',
'川 A9XT19' : '鄂 A9XT19',
'皖 AFB872' : '皖 AFR872',
'渝 DZ9968' : '晋 DZ9968',
'皖 A2N389' : '皖 AQN359',
'皖 E27765' : '皖 E97765',
'鲁 QA46D1' : '粤 QA4601',

'皖 AEZ152': '皖 AFZ152',
'皖 A88Z68': '皖 AS8Z68',
'闽 DEV572': '鄂 DEV572',
'皖 MR407': '皖 AMR407',
'鄂 A7172J': '贵 A7172J',
'皖 AHY612': '皖 ANY612',
'皖 AR8229': '皖 AK927W',
'皖 N2Q059': '皖 AMQ059',
'皖 AL666M': '皖 A1666M',
'皖 AZ15D2': '皖 AZ1502',
'苏 NT1199': '沪 NT1199',
'浙 F9K908': '赣 F9K908',
'皖 BZR331': '皖 RZR331',
'皖 AJ22J5': '皖 AH2W15',
'皖 ABH663': '皖 A8H467',
'鲁 CJ0001': '闽 CJ0001',
'皖 R0D963': '皖 R00963',
'粤 CVT910': '豫 CVT910',
'晋 BQ6M90': '浙 BQ6M90',
'皖 AD5R59': '皖 A05R59',
'皖 AX7M72': '皖 AX7M75',
'苏 E33P00': '苏 E33VC0',
'京 QXK818': '皖 QXK818',
'皖 AN7211': '皖 AN721L',
'浙 DK3785': '沪 DK3785',
'皖 AH7677': '皖 AH767W',
'皖 A166E4': '皖 A166F4',
'浙 ALC568': '青 ALC568',
'皖 AP1723': '皖 APH723',
'鲁 MEU969': '粤 HEU969',
'豫 ME4815': '赣 ME4815',
'皖 AFN286': '皖 APY862',
'皖 AZ6111': '皖 A76F11',
'鄂 A031JU': '豫 A031JU',
'鄂 AH448P': '粤 AH448P',
'浙 BK6857': '湘 BK6857',
'皖 A896P7': '皖 A896P1',
'皖 HR3241': '皖 NR3241',
'浙 LV0652': '津 LV0652',
'浙 F77X15': '渝 F77X15',
'浙 B886B8': '苏 B886B8',
'皖 AD5J86': '皖 AH5J86',
'皖 AJ0T64': '皖 AJ0T61',
'皖 AK0W20': '皖 AM0W20',
'皖 AXD162': '皖 AXD167',
'皖 ASL190': '皖 A9L190',
'浙 F368K1': '渝 F368K1',
'皖 AN9V96': '皖 AN9V16',
'皖 A97166': '皖 A97160',

'皖 A2R383' : '皖 A8R383',
'皖 A75D70' : '皖 A75H70',
'皖 H312V1' : '鲁 M712V1',
'川 A0791X' : '陕 A0791X',
'苏 K75F95' : '蒙 K75F95',
'皖 A23466' : '皖 A23464',
'苏 E23F07' : '苏 E23JP7',
'皖 AWP446' : '皖 AWF446',
'皖 AX8146' : '宁 AX8146',
'皖 HFG886' : '皖 N1G886',
'皖 AN7G5' : '皖 A787G5',
'皖 AD6383' : '皖 A06383',
'苏 E51Q07' : '苏 E51Q07',
'粤 A86QJ8' : '豫 A86QJ8',
'皖 AT560P' : '皖 AT560F',
'皖 AC6S66' : '皖 AC6S68',
'浙 A619V8' : '皖 A619V9',
'皖 ANZ536' : '皖 ANZ538',
'皖 A72JT9' : '皖 A72J15',
'鲁 RNX611' : '冀 RNX611',
'鄂 AT155V' : '豫 AT155V',
'皖 A83T7' : '皖 A83T07',
'川 AN5R29' : '冀 AN5R29',
'皖 A36031' : '皖 AEA031',
'闽 DHZ909' : '粤 DHZ909',
'皖 AB397G' : '皖 AB3978',
'皖 ACF6C1' : '皖 ACF601',
'皖 ASW324' : '皖 ASW824',
'皖 ATY735' : '皖 ATV735',
'皖 AX788V' : '皖 AX788W',
'浙 AR309R' : '鲁 AR309R',
'皖 A06D08' : '皖 A06DD8',
'粤 A892BU' : '陕 A892BU',
'皖 A42E80' : '皖 A42F80',
'皖 H5H606' : '皖 N5H606',
'鄂 AU70P5' : '陕 AU70P5',
'皖 AP0E86' : '皖 AP0L86',
'皖 AJ3684' : '皖 ABJ356',
'皖 ADU527' : '皖 A0U527',
'苏 B2711K' : '苏 B271UK',
'皖 AEE325' : '皖 AEL325',
'皖 AQN6C8' : '皖 AQN608',
'粤 BTW936' : '吉 BTW976',
'皖 AHX094' : '皖 AHXQ94',
'皖 ABSF79' : '皖 AD9F79',
'鲁 CXL543' : '粤 CXL543',
'浙 AZE739' : '浙 JZE739',
'皖 AWA6Q5' : '皖 AWA605',
'皖 A231Q7' : '皖 A23107',

'皖 AJH63Q' : '皖 AJH630',
'鄂 A85R83' : '赣 A85R83',
'皖 A22T96' : '皖 A22T99',
'皖 A6G6Z7' : '皖 A606Z7',
'皖 AYL867' : '皖 AY1867',
'皖 ADA5D5' : '皖 ADA505',
'川 A0K088' : '津 A0K088',
'皖 AQD975' : '皖 AQD915',
'皖 D4863J' : '皖 D48637',
'鄂 A6KC23' : '湘 A6KC23',
'皖 AK556M' : '黑 AK556M',
'浙 GD2500' : '苏 GD2500',
'鄂 AT770Z' : '川 AT770Z',
'赣 M26261' : '辽 D26261',
'皖 A80A35' : '闽 A80A35',
'皖 AHF024' : '皖 AHFQ24',
'粤 URE029' : '鲁 VRE029',
'皖 AT06QT' : '皖 AT060T',
'皖 AAU872' : '皖 AA0872',
'苏 AD876B' : '苏 JD876B',
'皖 AV1662' : '皖 AYU642',
'豫 NQ0551' : '京 NQ0551',
'鄂 A831SK' : '冀 A831SK',
'闽 BB12S9' : '粤 BB12S9',
'粤 B8Q82K' : '鲁 B8Q82K',
'鄂 A6E11K' : '川 A6E11K',
'浙 B3ZD36' : '苏 B3ZD36',
'皖 ANG812' : '皖 AM8812',
'鲁 Q19GN0' : '京 Q19BN0',
'皖 A97Z79' : '皖 AP7Z79',
'苏 AC822D' : '皖 AC822D',
'皖 A1F251' : '皖 A1E251',
'鄂 A9E03A' : '川 A9E03A',
'皖 A27CB3' : '皖 A27CD3',
'皖 AUE198' : '皖 AJE198',
'皖 AY918C' : '皖 AY915C',
'粤 BK18Q3' : '鲁 BK18Q3',
'皖 A9993' : '皖 ACH993',
'新 R8865A' : '豫 R8865A',
'皖 B547C0' : '皖 B54700',
'皖 AA621N' : '皖 AA631N',
'赣 E45042' : '鄂 E45042',
'浙 G6606N' : '浙 G6606K',
'皖 AJ262K' : '皖 AA262K',
'豫 MU2785' : '川 MU2785',
'皖 AH7279' : '苏 JH7279',
'皖 MH9E18' : '皖 HH9E18',
'赣 LBT526' : '蒙 LBT526',
'皖 ABS007' : '皖 ABS001',

```
'赣 L87A77': '辽 L87A77',
'粤 B0311C': '闽 B0311C',
'皖 A5R561': '皖 A5R551',
'浙 KJQ929': '蒙 KJQ929',
'京 MRY333': '津 MRY333',
'皖 AMX258': '鄂 AMX258',
'赣 L87283': '粤 L8T283',
'粤 B90KG9': '浙 B90KG9',
'皖 AD5Q8C': '皖 AD506C',
'川 AG8592': '京 AG8592',
'皖 AHD77G': '皖 AH077G',
'苏 AT58X8': '苏 AF58X8',
'鲁 MGE679': '苏 BE767W',
'皖 AF812Q': '皖 AFS120',
'皖 A3F619': '皖 ACF619',
'皖 A0Y0D0': '皖 A0Y000',
'皖 AUE118': '皖 AUF118',
'粤 S8Y585': '豫 S8Y585',
'川 EH4B78': '鄂 EH4P78',
'皖 A79565': '皖 A7S565',
'皖 AM062J': '晋 AM062J',
'皖 AG838U': '皖 AG838D',
'皖 AC7783': '皖 FC7783',
'赣 GVV921': '豫 GVV921',
'皖 AJ226': '皖 ADD226',
'京 N89N13': '豫 N89N13',
'皖 AE6719': '皖 AF6719',
'浙 DQ7222': '渝 DQ7222'}
```

As we can see, the model is most wrong on Chinese characters, I tried to fix it with augmentations, but still there are numbers on which the model is wrong. These errors most likely arise due to poor image quality, this can also be corrected by expanding the sample, or training the model specifically for recognizing Chinese characters and then merging it with the main model. Also, the model sometimes makes mistakes in the length of the number and in the middle characters. This can be fixed with the help of augmentations: for example, painting over some part of a certain symbol.

```
key_err_0 = list(errors.values())[0]
```

```
ind_err_0 = dataset_test.text.index(key_err_0)
```

Errors can also be due to damage to numbers, which we see in this example

```
torchvision.transforms.ToPILImage()(dataset_test[ind_err_0][0])
```



```
list(errors.keys())[0]
```

'皖 HL108D'

key_err_0

'皖 NL108D'

acc_avg = val_loop(dataloader_test, model, tokenizer, device)

Validation, acc: 0.9667

acc_avg

0.9666966696669667

Totals by metrics:

Accuracy on test: 0.9666966696669667

CER on test: 0.0058005800580058